

VE281

Data Structures and Algorithms

Open Addressing; Universal Hashing

Announcement

- Lecture Rescheduling
 - No course on Nov. 6, 8, 10, 13, 15.
 - Make-up courses on
 - Oct. 20, Oct. 27, Nov. 3, Nov. 24, Dec. 1 (Fridays): 2 pm – 3:40 pm.

Outline

- Collision Resolution: Open Addressing
 - Linear Probing
 - Quadratic Probing and Double Hashing
 - Performance of Open Addressing
- Pathological Data Sets and Universal Hashing

Open Addressing

- Reuse empty space in the hash table to hold colliding items.
- To do so, search the hash table in some systematic way for a bucket that is empty.
 - Idea: we use a sequence of hash functions h_0, h_1, h_2, \dots to probe the hash table until we find an empty slot.
 - I.e., we **probe** the hash table buckets mapped by $h_0(\text{key})$, $h_1(\text{key})$, \dots , in sequence, until we find an empty slot.
 - Generally, we could define $h_i(x) = h(x) + f(i)$

Open Addressing

- Three methods:

- Linear probing:

$$h_i(x) = (h(x) + i) \% n$$

- Quadratic probing:

$$h_i(x) = (h(x) + i^2) \% n$$

- Double hashing:

$$h_i(x) = (h(x) + i * g(x)) \% n$$

n is the hash table size

Linear Probing

$$h_i(\text{key}) = (h(\text{key}) + i) \% n$$

- Apply hash function h_0, h_1, \dots , in sequence until we find an empty slot.
 - This is equivalent to doing a linear search from $h(\text{key})$ until we find an empty slot.
- Example: Hash table size $n = 9$, $h(\text{key}) = \text{key} \% 9$
 - Thus $h_i(\text{key}) = (\text{key} \% 9 + i) \% 9$
 - Suppose we insert 1, 5, 11, 2, 17, 21, 31 in sequence

	1	11			5			
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]

How about 2?

Linear Probing

Example

- Hash table size $n = 9$, $h(\text{key}) = \text{key} \% 9$
 - Thus $h_i(\text{key}) = (\text{key} \% 9 + i) \% 9$
 - Suppose we insert 1, 5, 11, 2, 17, 21, 31 in sequence.

	1	11	2	21	5	31		17
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]

- $h_0(2) = 2$. Not empty!
- So we try $h_1(2) = 3$. It is empty, so we insert there!
- $h_0(21) = 3$. Not empty!
- $h_1(21) = 4$. It is empty, so we insert there!
- $h_0(31) = 4$. Not empty!
- $h_1(31) = 5$. Not empty!
- $h_2(31) = 6$. It is empty, so we insert there!

Linear Probing

find()

	1	11	2	21	5	31		17
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]

- With linear probing **$h_i(\text{key}) = (\text{key} \% 9 + i) \% 9$**
 - How will you **search** an item with key = 31?
 - How will you **search** an item with key = 10?
- Procedure: probe in the buckets given by $h_0(\text{key})$, $h_1(\text{key})$, ..., in sequence **until**
 - we find the key,
 - or we find an empty slot, which means the key is not found.

Linear Probing

remove()

	1	11	2	21	5	31		17
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]

- With linear probing $h_i(\text{key}) = (\text{key} \% 9 + i) \% 9$
 - How will you **remove** an item with key = 11?
 - If we just find 11 and delete it, will this work?

	1		2	21	5	31		17
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]

What is the result for searching key = 2 with the above hash table?

Linear Probing

remove()

cluster

	1		2	21	5	31		17
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]

- After deleting 11, we need to **rehash** the following “cluster” to fill the vacated bucket.
- However, we cannot move an item **beyond** its **actual** hash position. In this example, 5 cannot be moved ahead.

	1		2	21	5	31		17
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]

Linear Probing


Alternative implementation of remove()

	1	del	2	21	5	31		17
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]

- **Lazy deletion**: we mark deleted entry as “**deleted**”.
 - “deleted” is not the same as “empty”.
 - Now each bucket has three states: “occupied”, “empty”, and “deleted”.
- We can overwrite the “deleted” entry when inserting.
- When we **search**, we will keep looking if we encounter a “deleted” entry.

Linear Probing

Clustering Problem




- Clustering: when **contiguous** buckets are all occupied.
- **Claim**: Any hash value inside the cluster adds to **the end** of that cluster.
- Problems with a **large** cluster:
 - It becomes more likely that the next hash value will collide with the cluster.
 - Collisions in the cluster get more expensive to resolve.

Linear Probing

Clustering Problem

- Assuming input size N , table size $2N$:
 - What is the best-case cluster distribution?



- What is the worst-case cluster distribution?



- What's the average number of probes to find an empty slot for each case?

Outline

- Collision Resolution: Open Addressing
 - Linear Probing
 - Quadratic Probing and Double Hashing
 - Performance of Open Addressing
- Pathological Data Sets and Universal Hashing

Quadratic Probing

$$h_i(\text{key}) = (h(\text{key}) + i^2) \% n$$

- It is less likely to form large clusters.
- Example: Hash table size $n = 7$, $h(\text{key}) = \text{key} \% 7$
 - Thus $h_i(\text{key}) = (\text{key} \% 7 + i^2) \% 7$
 - Suppose we insert 9, 16, 11, 2 in sequence.

		9	16	11		2
[0]	[1]	[2]	[3]	[4]	[5]	[6]

- $h_0(16) = 2$. Not empty!
- $h_1(16) = 3$. It is empty, so we insert there.
- $h_0(2) = 2$. Not empty!
- $h_1(2) = 3$. Not empty!
- $h_2(2) = 6$. It is empty, so we insert there.

Problem of Quadratic Probing

- However, sometimes we will never find an empty slot even if the table isn't full!
- Luckily, if the **load factor** $L \leq 0.5$, we are guaranteed to find an empty slot.
 - Definition: given a hash table with n buckets that stores m objects, its **load factor** is

$$L = \frac{m}{n} = \frac{\text{\#objects in hash table}}{\text{\#buckets in hash table}}$$

More on Load Factor of Hash Table

- Question: which collision resolution strategy is feasible for load factor larger than 1?
 - Answer: separate chaining.
 - Note: for open addressing, we require $L \leq 1$.
- Claim: $L = O(1)$ is a necessary condition for operations to run in constant time.

Double Hashing

$$h_i(x) = (h(x) + i * g(x)) \% n$$

- Uses 2 distinct hash functions.
- Increment **differently** depending on the key.
 - If $h(x) = 13$, $g(x) = 17$, the probe sequence is 13, 30, 47, 64, ...
 - If $h(x) = 19$, $g(x) = 7$, the probe sequence is 19, 26, 33, 40, ...
 - For linear and quadratic probing, the incremental probing patterns are **the same** for all the keys.

Double Hashing

Example

- Hash table size $n = 7$, $h(\text{key}) = \text{key} \% 7$,
 $g(\text{key}) = (5 - \text{key}) \% 5$
 - Thus $h_i(\text{key}) = (\text{key} \% 7 + (5 - \text{key}) \% 5 * i) \% 7$
 - Suppose we insert 9, 16, 11, 2 in sequence.

		9		11	2	16
[0]	[1]	[2]	[3]	[4]	[5]	[6]

- $h_0(16) = 2$. Not empty!
- $h_1(16) = 6$. It is empty, so we insert there.
- $h_0(2) = 2$. Not empty!
- $h_1(2) = 5$. It is empty, so we insert there.

Outline

- Collision Resolution: Open Addressing
 - Linear Probing
 - Quadratic Probing and Double Hashing
 - Performance of Open Addressing
- Pathological Data Sets and Universal Hashing

Performance of Open Addressing

- Hard to analyze rigorously.
- The runtime is dominated by the number of comparisons.
- The number of comparisons depends on the load factor L .
- Define the expected number of comparisons in an **unsuccessful search** as $U(L)$.
- Define the expected number of comparisons in a **successful search** as $S(L)$.

Expected Number of Comparisons

- Linear probing

$$U(L) = \frac{1}{2} \left[1 + \left(\frac{1}{1-L} \right)^2 \right]$$
$$S(L) = \frac{1}{2} \left[1 + \frac{1}{1-L} \right]$$

L	$U(L)$	$S(L)$
0.5	2.5	1.5
0.75	8.5	2.5
0.9	50.5	5.5

$L \leq 0.75$ is recommended.

Expected Number of Comparisons

- Quadratic probing and double hashing

$$U(L) = \frac{1}{1 - L}$$
$$S(L) = \frac{1}{L} \ln \frac{1}{1 - L}$$

L	$U(L)$	$S(L)$
0.5	2	1.4
0.75	4	1.8
0.9	10	2.6

Which Strategy to Use?

- Both separate chaining and open addressing are used in real applications.
- Some basic guidelines:
 - If space is important, better to use open addressing.
 - If need removing items, better to use separate chaining.
 - **remove ()** is tricky in open addressing.
 - In mission critical application, prototype both and compare.

Outline

- Collision Resolution: Open Addressing
 - Linear Probing
 - Quadratic Probing and Double Hashing
 - Performance of Open Addressing
- Pathological Data Sets and Universal Hashing

Pathological Data Sets

- The **ideal** hash function spreads every data set out evenly.
- Does such **ideal** hash function exist?
 - No! For every hash function, there is a **pathological data set**.
- Reason: Fix a hash function $h: U \rightarrow \{0, 1, \dots, n - 1\}$
 - There exists a bucket i such that at least $|U|/n$ elements of U hash to i under h ...
 - ... if data set drawn only from these, everything collides!

Pathological Data Sets

- **Given**: A hash table with n buckets stores m items. Use separate chaining.
 - **Question**: If all m items are mapped to the same bucket, what's the time complexity of **find()**?
- **Answer**: The hash table degrades to a linked list.
 - The time complexity for **find()** is $O(m)$.
 - This is actually the **worst-case** time complexity.

Solution to Pathological Data Sets

- **Universal hashing:**
 - Design **a family** H of hash functions such that for **all** data set S , “**almost all**” functions $h \in H$ spread S out “**pretty evenly**”.
 - Pick a hash function **randomly** from the family H .

Universal Family of Hash Functions

- Definition: Let H be **a set of hash functions** from U to $\{0, 1, 2, \dots, n - 1\}$. H is **universal** if and only if:

- For all $x, y \in U$ with $x \neq y$,

$$\Pr_{h \in H}(h(x) = h(y)) \leq \frac{1}{n}$$

- In other words, any two keys of U collide with probability at most $1/n$ when the hash function h is chosen **uniformly at random** from H
- Note: keys x and y fixed. Random on hash function picked
 - At most $1/n$ of the total functions map x and y to the same bucket
- Collision probability is as small as our “gold standard” of **completely random hashing**

Universal Family of Hash Functions

Example

- Example #1: The set of **all** functions that map from U to $\{0, 1, 2, \dots, n - 1\}$.
 - The family contain $n^{|U|}$ functions.
- Is this family universal?
- Yes! Because for any keys $x \neq y$, exactly $1/n$ of the total functions map x and y to the same bucket
 - Partition all the functions into n^2 subsets $S_{i,j}$ ($0 \leq i, j \leq n - 1$), where $S_{i,j}$ contains functions h such that $h(x) = i$ and $h(y) = j$
 - The numbers of functions in all subsets are equal.

Universal Family of Hash Functions

Example

- Example #2: $\{h_0, h_1, \dots, h_{n-1}\}$ where $h_i: U \rightarrow i$, i.e., for any $u \in U$, $h_i(u) = i$.
- Is this family universal?
- No! Because for any keys $x \neq y$, all functions map x and y to the same bucket, i.e.,

$$\Pr_{h \in H}(h(x) = h(y)) = 1$$

Real Example: Hashing IP Addresses

- Let $U = \text{IP address of the form } (x_1, x_2, x_3, x_4)$ with each $x_i \in \{0, 1, \dots, 255\}$
- Let hash table size n be a **prime** number and $n > 255$.
 - Could be close to a multiple of #objects in the hash table.
- Define one hash function h_a per 4-tuple $a = (a_1, a_2, a_3, a_4)$ with each $a_i \in \{0, 1, \dots, n - 1\}$.
 - $h_a(x_1, x_2, x_3, x_4) = (a_1x_1 + a_2x_2 + a_3x_3 + a_4x_4) \bmod n$
 - There are n^4 such functions.

A Universal Family of Hash Functions

- **Define** the family $H =$ all n^4 h_a 's, i.e.,

$$H = \{h_a | a_1, a_2, a_3, a_4 \in \{0, 1, \dots, n-1\}\}$$

$$h_a(x_1, x_2, x_3, x_4) = (a_1x_1 + a_2x_2 + a_3x_3 + a_4x_4) \bmod n$$

Theorem: Family H is universal.

Proof

- Consider distinct IP addresses (x_1, x_2, x_3, x_4) and (y_1, y_2, y_3, y_4)
- Assume $x_4 \neq y_4$. We need to show the collision probability for a randomly chosen function $h_a \in H$ is at most $1/n$, i.e.,

$$\Pr_{h_a \in H} (h_a(x_1, \dots, x_4) = h_a(y_1, \dots, y_4)) \leq \frac{1}{n}$$

- Note: collision happens when

$$a_1x_1 + \dots + a_4x_4 = a_1y_1 + \dots + a_4y_4 \pmod{n}$$

$$\Leftrightarrow a_4(x_4 - y_4) = \sum_{i=1}^3 a_i(y_i - x_i) \pmod{n}$$

Next: let's fix random choice a_1, a_2, a_3 , but a_4 still random

Proof (cont.)

- Question: with a_1, a_2, a_3 fixed arbitrarily, how many choices of a_4 satisfy

$$a_4(x_4 - y_4) = \sum_{i=1}^3 a_i(y_i - x_i) \pmod{n} ?$$

- Note:

fixed value

1. $0 \leq x_4 \neq y_4 \leq 255$
 2. $n > 255$ is prime
- Claim: For any $b \in \{0, \dots, n-1\}$, there is at most one $a_4 \in \{0, \dots, n-1\}$ that let $a_4(x_4 - y_4) = b \pmod{n}$
 - Proof: for any $a_4' \in \{0, \dots, n-1\}$ and $a_4' \neq a_4$,
 $(a_4 - a_4')(x_4 - y_4) \neq 0 \pmod{n}$

$$\text{Imply } \Pr_{h_a \in H} (h_a(x_1, \dots, x_4) = h_a(y_1, \dots, y_4)) \leq \frac{1}{n}$$

Q.E.D.