

VE281

Data Structures and Algorithms

Shortest Path


Outline

- Shortest Path Problem
 - Unweighted Graph
 - Dijkstra's Algorithm
 - Bellman-Ford Algorithm

Shortest Path Problem

Introduction

- Given a weighted graph $G = (V, E)$, **path length** is defined as the sum of weights of edges on the path.
 - E.g., length of the path B, C, D, F is 12.
- Shortest path problem**: given a weighted graph $G = (V, E)$ and two nodes $s, d \in V$, find the shortest path from s to d .
 - Assume G is a directed graph without parallel edges of the same direction
 - For an undirected graph, we can replace each edge by two edges of the same weight but of different directions.



What is the shortest path from B to F?


Shortest Path Problem

- The starting node on the path is the **source** node and the ending node is the **destination** node.
- The previous problem is a **single source single destination** problem.
- What we will solve is a **single source all destinations** problem: Given $G = (V, E)$ and a node $s \in V$, find the shortest path from s to **every other** node in G .
 - **Single source single destination** problem can be solved by solving the **single source all destinations** problem.
 - However, **single source single destination** problem is not much easier than the **single source all destinations** problem.

Shortest Path Problem

A Simple Version: Unweighted Graphs

- For an unweighted graph, path length is defined as the number of edges on the path.
- How do you obtain the shortest path between a pair of nodes?




What is the shortest path from B to F?

Using breadth-first search!



Shortest Path for Unweighted Graphs

- Recall breadth-first search (BFS): Given a start node, visit all directly connected neighbors first, then nodes 2 hops away, 3 hops away, and so on.
 - This is exactly what we want!
 - When the node visited is the destination node, we stop.
 - When the queue becomes empty, there is no path between the two nodes.



Shortest Path for Unweighted Graphs

- Additional bookkeeping
 - Store the distance.
 - Store the **predecessor** on the shortest path, i.e., the previous node on the path.





	A	B	C	D	E	F
dist	1	0	1	2	2	3
pred	B	-	B	A	C	D

Shortest Path for Unweighted Graphs


- We can obtain the shortest path by backtracking.

- E.g., shortest path from B to F

B → A → D → F



	A	B	C	D	E	F
dist	1	0	1	2	2	3
prev	B	-	B	A	C	D




Outline

- Shortest Path Problem
 - Unweighted Graph
 - Dijkstra's Algorithm
 - Bellman-Ford Algorithm

Shortest Path for Weighted Graphs

- The problem becomes more difficult when edges have different weights.

- Breadth-first search won't work!
- What is the shortest path from B to F?



- If the weights are **non-negative**, then we can apply **Dijkstra's Algorithm** (more details & examples from Ve203)
 - Works only when all weights are non-negative
 - A greedy algorithm for solving single source all destinations shortest path problem

Dijkstra's Algorithm

- Keep **distance estimate** $D(v)$ and **predecessor** $P(v)$ for each node v .
 - Predecessor: the previous node on the shortest path.
- 1. Initially, $D(s) = 0$; $D(v)$ for other nodes is $+\infty$; $P(v)$ is unknown.
- 2. Store all the nodes in a set R .
- 3. While R is not empty
 - 1. Choose node v in R such that $D(v)$ is the **smallest**. Remove v from the set R .
 - 2. Declare that v 's shortest distance is known, which is $D(v)$.
 - 3. For each of v 's **neighbors** u that is **still in R** , update distance estimate $D(u)$ and predecessor $P(u)$.

Updating

- For each of v 's **neighbors** u that is **still in R** , if $D(v) + w(v, u) < D(u)$, then update $D(u) = D(v) + w(v, u)$ and the predecessor $P(u) = v$.
- I.e., update $D(u)$ if the path going through v is shorter than the best path so far to u .

Dijkstra's Algorithm v.s. Prim's Algorithm

- Dijkstra's algorithm is similar to Prim's algorithm
 - Prim's algorithm: grow the set of nodes we add to the MST.
 - Dijkstra's algorithm: grow the set of nodes to which we know the shortest path.

Dijkstra's Algorithm

Time Complexity

- Number of times to find the smallest $D(v)$: $|V|$.
 - Each cost? Linear scan: $O(|V|)$; Binary heap: $O(\log |V|)$; Fibonacci heap: $O(\log |V|)$
- Total number of times to update the neighbors: $|E|$.
 - Since each neighbor of each node could be potentially updated.
 - Each cost? Linear scan: $O(1)$; Binary heap: $O(\log |V|)$; Fibonacci heap: $O(1)$
- Total time complexity
 - Linear scan: $O(|E| + |V|^2) = O(|V|^2)$
 - Binary heap: $O(|V| \log |V| + |E| \log |V|)$
 - Fibonacci heap: $O(|V| \log |V| + |E|)$


Outline

- Shortest Path Problem
 - Unweighted Graph
 - Dijkstra's Algorithm
 - Bellman-Ford Algorithm

Limitation of Dijkstra's Algorithm

- Not always correct on graphs with negative edge weights

Find shortest paths
from source node A




- Solution: The Bellman-Ford algorithm

Negative Cycles

- If the graph contains a negative cycle reachable from the source S , then the shortest paths for some destinations are
— ∞

- Can traverse the negative cycle infinite times
- Thus, if there is any negative cycle reachable from S , the short path problem is not well defined



- If there is no negative cycle reachable from S , Bellman-Ford algorithm reports the correct shortest paths for all destinations
- If there is such a negative cycle, the algorithm will report it

Key Idea of Bellman-Ford Algorithm


- Claim: Suppose there is no negative cycles. Then, for each node $v \in V$, there is a shortest $s-v$ path with $\leq |V| - 1$ edges
 - Proof: If shortest path has $\geq |V|$ edges, then there is a cycle. By assumption, the cycle has non-negative length. We can drop the cycle and obtain a no-worse path
- Key idea: consider shortest $s-v$ path with **at most** i edges
 - Denoted as $P(v, i)$
 - It can be recursively obtained (show next)
 - If there is no negative cycles, then $P(v, |V| - 1)$ is a shortest $s-v$ path

Optimality

- Define $P(v, i)$ as the shortest s - v path with **at most** i edges
- Case 1: $P(v, i)$ has $\leq i - 1$ edges. Then $P(v, i)$ is also a shortest s - v path with $\leq i - 1$ edges.
 - Proof: By contradiction, assume Q is an s - v path with $\leq i - 1$ edges shorter than $P(v, i)$. Then $P(v, i)$ cannot be the shortest s - v path with $\leq i$ edges

Optimality

- Case 2: $P(v, i)$ has i edges. Suppose the first $i - 1$ edges of $P(v, i)$ form the path P' and the last edge is (u, v) , then P' is a shortest s - u path with $\leq i - 1$ edges.
- Proof: By contradiction, suppose Q is a s - u path with $\leq i - 1$ edges shorter than P' . Then $Q + (u, v)$ is shorter than $P' + (u, v)$. Then $P(v, i)$ cannot be the shortest s - v path with $\leq i$ edges



Recurrence


- Possible candidates for $P(v, i)$:
 - $P(v, i - 1)$
 - $P(u, i - 1) + (u, v)$, for all $(u, v) \in E$
- $P(v, i)$ is the shortest one of the above candidates
- Define the length of $P(v, i)$ as $L(v, i)$
- $$L(v, i) = \min\{L(v, i - 1), \min_{(u,v) \in E} [L(u, i - 1) + w(u, v)]\}$$

The Bellman-Ford Algorithm




- Initially, $L(s, 0) = 0$; $L(v, 0) = +\infty$ for other $v \in V$
- For $i = 1, 2, \dots, |V| - 1$
 - For each $v \in V$,
$$L(v, i) = \min\{L(v, i - 1), \min_{(u,v) \in E} [L(u, i - 1) + w(u, v)]\}$$
- If there is no negative cycle, the length of shortest s - v path is $L(v, |V| - 1)$ for all $v \in V$
- Can keep the predecessor to track shortest path

Example

- Let start point be A




Initialization




Example


$i = 2$



$i = 3$



$i = 4$



Final shortest path
length

Time Complexity of Bellman-Ford

- Initially, $L(s, 0) = 0$; $L(v, 0) = +\infty$ for other $v \in V$
- For $i = 1, 2, \dots, |V| - 1$
 - For each $v \in V$,
$$L(v, i) = \min\{L(v, i - 1), \min_{(u,v) \in E} [L(u, i - 1) + w(u, v)]\}$$
- Assume adjacency list representation. What's the time complexity?
 - $O((|V| + |E|)|V|)$ (worse than Dijkstra's algorithm)

Detecting Negative Cycle

- We can modify Bellman-Ford algorithm to report the existence of negative cycle by running it for one more iteration to obtain $L(v, |V|)$ for all $v \in V$
- Claim: G has no negative cycle reachable from $s \Leftrightarrow L(v, |V| - 1) = L(v, |V|)$ for all $v \in V$
- Proof of \Rightarrow : There exists a shortest s - v path with $\leq |V| - 1$ edges. Thus $L(v, |V| - 1) = L(v, |V|)$ for all $v \in V$

Detecting Negative Cycle

- Claim: G has no negative cycle reachable from $s \Leftrightarrow L(v, |V| - 1) = L(v, |V|)$ for all $v \in V$
- Proof of \Leftarrow :
 - Let $d(v) = L(v, |V| - 1)(= L(v, |V|))$
 - By the recurrence
$$\underline{L(v, |V|)} = \min \left\{ L(v, |V| - 1), \min_{(u,v) \in E} \left[\underline{L(u, |V| - 1)} + w(u, v) \right] \right\},$$

$d(v)$ $d(u)$
 - ... we have for any $(u, v) \in E$, $d(v) \leq d(u) + w(u, v)$, or $d(v) - d(u) \leq w(u, v)$
 - Now consider any cycle C , we have $\sum_{(u,v) \in C} w(u, v) \geq 0$